

# Introduction to Linux

A practical guide for competitive programmers

# Agenda

## Fundamentals

1. **Introduction** — What is Linux, distributions, why CP
2. **Getting Started** — WSL, setup, resources
3. **Terminal Basics** — Tips, wildcards, man, arguments
4. **Filesystem** — Directory structure, paths, nano


## Practical Skills

5. **Essential Commands** — ls, cd, cp, mv, rm, cat, ...
6. **Permissions & System** — chmod, sudo, apt
7. **C++ Compilation** — g++, flags, warnings, Makefile
8. **I/O Redirection** — Redirects, piping, chaining

Interactive playground available at [linux-playground.vercel.app](https://linux-playground.vercel.app)

# What is Linux?

- A **free, open-source** operating system — like Windows or macOS
- Built on the **Linux kernel**, created by Linus Torvalds in 1991
- Powers most servers, supercomputers, and Android devices
- Highly customizable — you can inspect and modify anything



# Linux Distributions

A **distribution** bundles the Linux kernel with software, package managers, and desktop environments.

## Ubuntu

Most popular desktop distro. Great GUI & community.

## Red Hat

Enterprise-focused.  
Corporate servers.

## Fedora

Cutting-edge features.  
Community-driven.

## Arch Linux

Minimalist. Highly customizable.

# Linux Distributions

A **distribution** bundles the Linux kernel with software, package managers, and desktop environments.

## Ubuntu

Most popular desktop distro. Great GUI & community.

## Red Hat

Enterprise-focused.  
Corporate servers.

## Fedora

Cutting-edge features.  
Community-driven.

## Arch Linux

Minimalist. Highly customizable.

All distros share the same kernel and shell commands – skills transfer across distributions.

# Why Linux for Competitive Programming?

# Why Linux for Competitive Programming?

- **Competition environments use Linux** – NOI, TFT, and many ICPC regionals run on Linux

# Why Linux for Competitive Programming?

- **Competition environments use Linux** – NOI, TFT, and many ICPC regionals run on Linux
- **Free and open source** – no licensing costs, install on any machine

# Why Linux for Competitive Programming?

- **Competition environments use Linux** – NOI, TFT, and many ICPC regionals run on Linux
- **Free and open source** – no licensing costs, install on any machine
- **Superior shell** – bash/zsh syntax is far more powerful than Windows cmd

# Why Linux for Competitive Programming?

- **Competition environments use Linux** – NOI, TFT, and many ICPC regionals run on Linux
- **Free and open source** – no licensing costs, install on any machine
- **Superior shell** – bash/zsh syntax is far more powerful than Windows cmd
- **Portable scripts** – shell commands work across all Linux distributions

# Why Linux for Competitive Programming?

- **Competition environments use Linux** – NOI, TFT, and many ICPC regionals run on Linux
- **Free and open source** – no licensing costs, install on any machine
- **Superior shell** – bash/zsh syntax is far more powerful than Windows cmd
- **Portable scripts** – shell commands work across all Linux distributions
- **Developer ecosystem** – g++, gdb, valgrind, and other tools are first-class citizens

# Which Distribution Should I Use?

## Ubuntu

- Excellent GUI out of the box
- Largest community & docs
- Best for beginners

**Get it:** [ubuntu.com](http://ubuntu.com)

## WSL (Windows)

- Linux shell inside Windows
- No VM needed, great performance
- Access Windows files from Linux

**Get it:** Built into Windows  
10/11

## Arch Linux

- Extremely lightweight
- Highly customizable
- Learn Linux internals

**Get it:** [archlinux.org](http://archlinux.org)

**Recommendation:** Use WSL if you're on Windows. Use Ubuntu if you want a dedicated Linux machine.

# Getting Started

WSL, setup, and resources

# Windows Subsystem for Linux (WSL)

WSL lets you run a full Linux environment on Windows — no VM or dual boot needed.

## WSL 1

- Translates Linux syscalls to Windows
- **Faster** cross-OS filesystem access
- No full Linux kernel
- Limited hardware support

## WSL 2

- Runs a **real Linux kernel**
- Full system call compatibility
- **Docker** and **CUDA** support
- Slightly slower cross-OS file access

# Windows Subsystem for Linux (WSL)

WSL lets you run a full Linux environment on Windows — no VM or dual boot needed.

## WSL 1

- Translates Linux syscalls to Windows
- **Faster** cross-OS filesystem access
- No full Linux kernel
- Limited hardware support

## WSL 2

- Runs a **real Linux kernel**
- Full system call compatibility
- **Docker** and **CUDA** support
- Slightly slower cross-OS file access

For competitive programming, WSL 2 is recommended — full kernel means everything just works.

# Setting Up WSL

# Setting Up WSL

1. Open **PowerShell as Administrator** and run:

```
wsl --install
```

# Setting Up WSL

1. Open **PowerShell as Administrator** and run:

```
wsl --install
```

2. **Restart** your computer when prompted

# Setting Up WSL

1. Open **PowerShell as Administrator** and run:

```
wsl --install
```

2. **Restart** your computer when prompted
3. Open **Ubuntu** from the Start menu — first launch completes setup

# Setting Up WSL

1. Open **PowerShell as Administrator** and run:

```
wsl --install
```

2. **Restart** your computer when prompted
3. Open **Ubuntu** from the Start menu — first launch completes setup
4. Create a **username and password** when prompted

# Setting Up WSL

1. Open **PowerShell as Administrator** and run:

```
wsl --install
```

2. **Restart** your computer when prompted
3. Open **Ubuntu** from the Start menu — first launch completes setup
4. Create a **username and password** when prompted
5. Update your system:

```
sudo apt update && sudo apt upgrade -y
```

# Setting Up WSL

1. Open **PowerShell as Administrator** and run:

```
wsl --install
```

2. **Restart** your computer when prompted
3. Open **Ubuntu** from the Start menu — first launch completes setup
4. Create a **username and password** when prompted
5. Update your system:

```
sudo apt update && sudo apt upgrade -y
```

**Requirements:** Windows 10 version 2004+ or any Windows 11. Run `winver` to check.

# How to Learn Linux Efficiently

## Resources

- **Google it** – most problems have been solved
- **Stack Overflow** – Q&A for specific issues
- **ChatGPT / AI** – great for explaining commands
- **man pages** – built-in manuals for every command
- **This presentation** – bookmark for reference

## Practice

- **Use Linux daily** – muscle memory matters
- **Try the Linux Playground** – interactive missions
- **Break things** – that's how you learn (use WSL)
- **Read error messages** – they tell you what's wrong

# Linux Playground – Interactive Practice

Try commands in your browser with guided missions:

The screenshot shows the Linux Playground interface. On the left, a sidebar displays the title "Linux Playground" with a progress bar indicating "0/6 Levels Complete". Below this, a mission titled "Project Scaffolding" (L1) is shown with a completion status of "0/5". The mission description reads: "Organize your workspace. A real competitive programmer keeps a clean directory tree. You must use the correct flags — no lazy shortcuts." A numbered list of tasks is provided, starting with "1. Create the nested path CP/Codeforces/Round900 using a single mkdir -p command." A "Hint" button is also present. To the right, a large workspace area displays a welcome message and instructions: "Welcome to Linux Playground – Competitive Programming Edition. Type 'help' for available commands. Complete the missions on the left panel." A terminal prompt "user@linux:~\$" is visible at the bottom of the workspace.

Welcome to Linux Playground – Competitive Programming Edition

Type "help" for available commands. Complete the missions on the left panel.

user@linux:~\$

Visit [linux-playground.vercel.app](https://linux-playground.vercel.app) – 6 mission levels from basic navigation to contest simulation.

# Terminal Basics

Tips, wildcards, manual, and arguments

# Opening the Terminal

## Ubuntu

Press `Ctrl` + `Alt` + `T` to open the default terminal.

Or search for "Terminal" in the application menu.

## WSL

- Open **Windows Terminal** → select Ubuntu tab
- Or type `wsl` in PowerShell
- Or search "Ubuntu" in the Start menu

## What You'll See

```
{ user@hostname:~$
```

This is your **shell prompt** – it shows the username, machine name, and current directory (`~` = home).

# Terminal Tips & Shortcuts

# Terminal Tips & Shortcuts

- **Tab completion** — type a few characters, press `Tab` to auto-complete
  - Press `Tab` twice to list all matching candidates

# Terminal Tips & Shortcuts

- **Tab completion** – type a few characters, press `Tab` to auto-complete
  - Press `Tab` twice to list all matching candidates
- **Command history** – `↑` / `↓` to cycle through previous commands

# Terminal Tips & Shortcuts

- **Tab completion** — type a few characters, press `Tab` to auto-complete
  - Press `Tab` twice to list all matching candidates
- **Command history** — `↑` / `↓` to cycle through previous commands
- **Cancel input** — `Ctrl` + `C` to discard the current line

# Terminal Tips & Shortcuts

- **Tab completion** – type a few characters, press `Tab` to auto-complete
  - Press `Tab` twice to list all matching candidates
- **Command history** – `↑` / `↓` to cycle through previous commands
- **Cancel input** – `Ctrl + C` to discard the current line
- **Clear screen** – `clear` or `Ctrl + L`

# Terminal Tips & Shortcuts

- **Tab completion** – type a few characters, press `Tab` to auto-complete
  - Press `Tab` twice to list all matching candidates
- **Command history** – `↑` / `↓` to cycle through previous commands
- **Cancel input** – `Ctrl + C` to discard the current line
- **Clear screen** – `clear` or `Ctrl + L`
- **Move cursor** – `Ctrl + A` (start), `Ctrl + E` (end of line)

# Terminal Tips & Shortcuts

- **Tab completion** – type a few characters, press `Tab` to auto-complete
  - Press `Tab` twice to list all matching candidates
- **Command history** – `↑` / `↓` to cycle through previous commands
- **Cancel input** – `Ctrl + C` to discard the current line
- **Clear screen** – `clear` or `Ctrl + L`
- **Move cursor** – `Ctrl + A` (start), `Ctrl + E` (end of line)

**Example:** To reference `this_is_a_super_long_filename`, type `this_is` then press `Tab` – the terminal completes the rest.

# Wildcard Character – \*

The \* matches **any sequence of characters** in file/directory names.

```
# Directory contains: 1.txt, 2.txt, 3.txt, notes.md
rm *.txt          # Removes 1.txt, 2.txt, 3.txt (not notes.md)
ls *.md           # Lists only notes.md
```

# Wildcard Character – \*

The \* matches **any sequence of characters** in file/directory names.

```
# Directory contains: 1.txt, 2.txt, 3.txt, notes.md  
rm *.txt          # Removes 1.txt, 2.txt, 3.txt (not notes.md)  
ls *.md           # Lists only notes.md
```

More examples with files `file123.txt` and `file456.txt` :

```
ls file*.txt      # Matches both - file + anything + .txt  
ls file*          # Matches both - file + anything  
ls f*.txt         # Matches both - f + anything + .txt
```

# Wildcard Character – \*

The \* matches **any sequence of characters** in file/directory names.

```
# Directory contains: 1.txt, 2.txt, 3.txt, notes.md  
rm *.txt          # Removes 1.txt, 2.txt, 3.txt (not notes.md)  
ls *.md           # Lists only notes.md
```

More examples with files `file123.txt` and `file456.txt` :

```
ls file*.txt      # Matches both – file + anything + .txt  
ls file*          # Matches both – file + anything  
ls f*.txt         # Matches both – f + anything + .txt
```

**Warning:** `rm *` deletes **everything** in the current directory. Always verify with `ls` first!

# Reading the Manual — `man`

Every command has a built-in manual page:

```
man cp      # Manual for the cp (copy) command  
man ls      # Manual for the ls (list) command  
man zip     # Manual for the zip command
```

# Reading the Manual — `man`

Every command has a built-in manual page:

```
man cp      # Manual for the cp (copy) command  
man ls      # Manual for the ls (list) command  
man zip     # Manual for the zip command
```

## Navigating man pages

Key	Action
↑ / ↓	Scroll line by line
Space	Scroll one page down
/pattern	Search for text
q	Quit

# Reading the Manual — `man`

Every command has a built-in manual page:

```
man cp      # Manual for the cp (copy) command  
man ls      # Manual for the ls (list) command  
man zip     # Manual for the zip command
```

## Navigating man pages

Key	Action
↑ / ↓	Scroll line by line
Space	Scroll one page down
/pattern	Search for text
q	Quit

**Quick alternative:** Most commands support `-h` or `--help` for a shorter summary.

# Command Arguments

Arguments modify how a command behaves – like function parameters in code.

Short form (single `-`)

```
ls -a          # all files  
ls -l          # long format  
ls -la         # combined  
rm -rf folder/ # recursive + force
```

Long form (double `--`)

```
ls --all  
rm --recursive --force folder/  
g++ --version
```

# Command Arguments

Arguments modify how a command behaves — like function parameters in code.

Short form (single `-`)

```
ls -a          # all files  
ls -l          # long format  
ls -la         # combined  
rm -rf folder/ # recursive + force
```

Long form (double `--`)

```
ls --all  
rm --recursive --force folder/  
g++ --version
```

Common flag conventions

`-f` **force**

`-v` **verbose / version**

`-a` **all**

`-r` **recursive**

`-i` **interactive**

`-o` **output**

`-h` **help**


Check the manual or `--help` for any command's full list of flags.

# Linux Filesystem

Directory structure, paths, and navigation

# Filesystem Overview

Linux organizes everything in a single tree rooted at `/`.



Unlike Windows (`C:\`, `D:\`), Linux has a **single root** `/` — everything is a subdirectory, including external drives.

# Special Directories

Symbol	Name	Meaning
/	Root	Top of entire filesystem
~	Home	User's home ( <code>/home/user</code> )
.	Current	Directory you're in now
..	Parent	One level up

```
# Current directory: /home/user/projects
pwd                  # /home/user/projects
cd ..                # Now at /home/user
cd ~                 # Now at /home/user
cd /                 # Now at root
cd ~/projects        # Back to /home/user/projects
```

# Special Directories

Symbol	Name	Meaning
/	Root	Top of entire filesystem
~	Home	User's home ( <code>/home/user</code> )
.	Current	Directory you're in now
..	Parent	One level up

```
# Current directory: /home/user/projects
pwd                  # /home/user/projects
cd ..                # Now at /home/user
cd ~                 # Now at /home/user
cd /                 # Now at root
cd ~/projects        # Back to /home/user/projects
```

Your terminal always starts at `~` (home directory) unless you change it with `cd`.

# Key System Directories

`/bin` — Essential command executables: `ls`,  
`cp`, `mv`, `cat`

`/dev` — Device files. `/dev/null` discards all data written to it (a "black hole")

`/home` — User directories: `/home/user1`,  
`/home/user2`, etc.

`/lib` — Shared libraries needed by `/bin` and system binaries

`/mnt` — Where external drives and filesystems are temporarily mounted

`/etc` — System-wide config files (network, users, services)

# Relative vs Absolute Paths

## Absolute path

Starts from root `/` — always the same regardless of where you are.

```
{ /home/user/projects/solution.cpp }
```

## Relative path

Relative to your current directory.

```
{ ./projects/solution.cpp  
  ../user2/file.txt }
```

## Example from `/home/user`

Target	Absolute	Relative
File in projects	/home/user/projects/a.cpp	./projects/a.cpp
Parent	/home	..
Other user	/home/user2/file.txt	../user2/file.txt

# Text Editor in Terminal – nano

`nano` is a simple, beginner-friendly terminal text editor.

```
{ nano myfile.cpp      # Open (or create) myfile.cpp for editing
```

## Essential shortcuts

Shortcut	Action	Shortcut	Action
<code>Ctrl + X</code>	Exit nano	<code>Ctrl + O</code>	Save without exiting
<code>Y → Enter</code>	Save on exit	<code>Ctrl + K</code>	Cut current line
<code>N → Enter</code>	Discard on exit	<code>Ctrl + W</code>	Search text

For advanced editing, look into **vim** or **VS Code** with remote SSH.

# Essential Commands

File and directory operations

# Command Overview

## Navigation

- `pwd` — print working directory
- `ls` — list directory contents
- `cd` — change directory

## File Operations

- `mkdir` — create directories
- `touch` — create empty files
- `cp` — copy files/directories
- `mv` — move or rename
- `rm` — remove files/directories

## Content & Search

- `cat` — display file contents
- `echo` — print / write to files
- `diff` — compare files
- `grep` — search file contents
- `zip` / `unzip` — archive files

# List Directory Contents – `ls`

```
{ ls          # List files in current directory  
  ls -a      # Include hidden files (names starting with .)  
  ls -l      # Long format – permissions, size, date  
  ls -la     # Both: hidden files + long format  
  ls /some/path # List a specific directory
```

# List Directory Contents – `ls`

```
{ ls          # List files in current directory  
  ls -a      # Include hidden files (names starting with .)  
  ls -l      # Long format – permissions, size, date  
  ls -la     # Both: hidden files + long format  
  ls /some/path # List a specific directory
```

Example output of `ls -la`

```
{ drwxr-xr-x  4 user user 4096 Feb 15 10:30 .  
drwxr-xr-x  3 user user 4096 Feb 14 09:00 ..  
-rw-r--r--  1 user user  220 Feb 14 09:00 .bashrc  
drwxr-xr-x  2 user user 4096 Feb 15 10:30 projects  
-rw-r--r--  1 user user   128 Feb 15 10:25 solution.cpp
```

`d` = directory, `-` = file. Hidden files start with `.` (only shown with `-a`).

# Change Directory — `cd`

```
{ cd /          # Go to root directory  
cd ~          # Go to home directory (/home/user)  
cd ..         # Go up one level (parent directory)  
cd projects   # Enter the "projects" subdirectory  
cd ~/projects # Go to projects under home (absolute via ~)
```

# Change Directory — `cd`

```
{ cd /          # Go to root directory  
  cd ~         # Go to home directory (~ /home/user)  
  cd ..        # Go up one level (parent directory)  
  cd projects  # Enter the "projects" subdirectory  
  cd ~/projects # Go to projects under home (absolute via ~)
```

## Practical navigation

```
{ user@linux:~$ cd projects/contest  
user@linux:~/projects/contest$ ls  
A.cpp  B.cpp  input.txt  
user@linux:~/projects/contest$ cd ../../  
user@linux:~$ pwd  
/home/user
```

The prompt updates to show your current path — always know where you are.

# Create Directories — `mkdir`

```
{ mkdir mydir          # Create a single directory  
  mkdir dir1 dir2 dir3    # Create multiple directories  
  mkdir -p CP/Codeforces/Round900  # Create nested directories (parents too)
```

# Create Directories — `mkdir`

```
{ mkdir mydir          # Create a single directory  
  mkdir dir1 dir2 dir3 # Create multiple directories  
  mkdir -p CP/Codeforces/Round900 # Create nested directories (parents too)
```

Why `-p` ?

Without it, creating `CP/Codeforces/Round900` fails if parent dirs don't exist.

```
{ mkdir CP/Codeforces/Round900      # Error: No such file or directory  
  mkdir -p CP/Codeforces/Round900 # Creates all missing parents automatically
```

# Create Files – touch

```
{ touch file.txt          # Create an empty file (or update timestamp)
  touch a.cpp b.cpp c.cpp # Create multiple files at once
```

# Create Files – touch

```
{ touch file.txt          # Create an empty file (or update timestamp)  
  touch a.cpp b.cpp c.cpp # Create multiple files at once
```

Alternative: create with >

```
{ > file.txt            # Also creates an empty file (overwrites if exists)  
  echo "content" > file.txt # Create a file with content
```

# Create Files – `touch`

```
{ touch file.txt          # Create an empty file (or update timestamp)  
  touch a.cpp b.cpp c.cpp # Create multiple files at once
```

Alternative: create with `>`

```
{ > file.txt            # Also creates an empty file (overwrites if exists)  
  echo "content" > file.txt # Create a file with content
```

`touch` is safe on existing files – it only updates the timestamp without changing content. `> file.txt` will **erase** an existing file's content.

# Copy – cp

```
{ cp source.txt dest.txt          # Copy a file  
  cp file1.txt file2.txt dest_folder/    # Copy multiple files into a directory  
  cp -r source_dir/ dest_dir/        # Copy a directory recursively
```

# Copy – cp

```
{ cp source.txt dest.txt          # Copy a file  
  cp file1.txt file2.txt dest_folder/    # Copy multiple files into a directory  
  cp -r source_dir/ dest_dir/        # Copy a directory recursively
```

## Key points

- **File** copy – no flags needed
- **Directory** copy – requires `-r` (recursive), otherwise error
- Overwrites existing destination files silently (`-i` to prompt, `-f` to force)

```
{ cp solution.cpp backup_solution.cpp      # Quick backup  
  cp -r contest/ contest_backup/        # Backup entire directory
```

# Move / Rename — mv

```
{ mv source.txt dest_folder/          # Move file into a directory  
  mv old_name.cpp new_name.cpp      # Rename a file  
  mv dir1/ dir2/                   # Move (or rename) a directory
```

# Move / Rename — `mv`

```
{ mv source.txt dest_folder/          # Move file into a directory  
  mv old_name.cpp new_name.cpp      # Rename a file  
  mv dir1/ dir2/                   # Move (or rename) a directory
```

`mv` vs `cp`

`cp`

`mv`

Original file

Kept

Removed

Directory flag

Needs `-r`

No flag needed

Use case

Duplicate

Relocate or rename

`mv` is the standard way to **rename** files and directories in Linux.



# Remove — `rm`

```
{ rm file.txt          # Remove a single file  
  rm file1 file2 file3 # Remove multiple files  
  rm -r directory/    # Remove a directory and all its contents  
  rm -rf directory/   # Force remove – no prompts, even read-only files
```

`rm` is permanent — there is no recycle bin. `rm -rf /` would delete your entire filesystem. Always double-check paths and verify with `ls` first!

# Remove — `rm`

```
{ rm file.txt           # Remove a single file  
  rm file1 file2 file3 # Remove multiple files  
  rm -r directory/    # Remove a directory and all its contents  
  rm -rf directory/   # Force remove – no prompts, even read-only files
```

`rm` is permanent — there is no recycle bin. `rm -rf /` would delete your entire filesystem. Always double-check paths and verify with `ls` first!

```
# Safe workflow: verify before deleting  
ls *.o                  # Check which .o files exist  
rm *.o                  # Now remove them
```

# Display File Contents – `cat`

```
{ cat file.txt          # Print file contents to terminal  
  cat file1.txt file2.txt # Print multiple files in sequence (concatenate)
```

# Display File Contents – `cat`

```
{ cat file.txt          # Print file contents to terminal  
  cat file1.txt file2.txt # Print multiple files in sequence (concatenate)
```

## Example

```
{ user@linux:~$ cat hello.cpp  
#include <bits/stdc++.h>  
using namespace std;  
int main() {  
    cout << "Hello, World!" << endl;  
}
```

Useful for quickly viewing short files, checking I/O files, or piping content to other commands.

# Display File Contents – `cat`

```
{ cat file.txt          # Print file contents to terminal  
  cat file1.txt file2.txt # Print multiple files in sequence (concatenate)
```

## Example

```
{ user@linux:~$ cat hello.cpp  
#include <bits/stdc++.h>  
using namespace std;  
int main() {  
    cout << "Hello, World!" << endl;  
}
```

Useful for quickly viewing short files, checking I/O files, or piping content to other commands.

For long files, use `less file.txt` (scrollable) or `head -n 20 file.txt` (first 20 lines).

# Print & Write – echo

```
echo "Hello, World!"          # Print text to terminal  
echo "Hello" > file.txt      # Write to file (overwrites)  
echo "More text" >> file.txt # Append to file
```

# Print & Write – echo

```
echo "Hello, World!"          # Print text to terminal  
echo "Hello" > file.txt       # Write to file (overwrites)  
echo "More text" >> file.txt # Append to file
```

> VS >>

## Operator

>

## Behavior

**Overwrites** the file (creates if doesn't exist)

>>

**Appends** to the end of the file

```
echo "line 1" > output.txt    # output.txt → "line 1"  
echo "line 2" > output.txt    # output.txt → "line 2" (line 1 gone!)  
echo "line 3" >> output.txt  # output.txt → "line 2\nline 3"
```

# Compare Files — `diff`

```
{ diff file1.txt file2.txt      # Compare two files line by line  
    diff -Z a.out b.out        # Ignore trailing whitespace
```

# Compare Files — `diff`

```
{ diff file1.txt file2.txt      # Compare two files line by line  
  diff -Z a.out b.out        # Ignore trailing whitespace
```

## Reading diff output

```
{ user@linux:~$ diff expected.out my.out  
3c3  
< 42  
---  
> 41
```

**3c3** = line 3 changed · < = first file · > = second file · **No output** = identical

# Compare Files — `diff`

```
{ diff file1.txt file2.txt      # Compare two files line by line  
  diff -Z a.out b.out        # Ignore trailing whitespace
```

## Reading diff output

```
{ user@linux:~$ diff expected.out my.out  
3c3  
< 42  
---  
> 41
```

**3c3** = line 3 changed · < = first file · > = second file · **No output** = identical

**CP tip:** Use `diff -Z` to ignore trailing spaces — matches how most online judges compare output.

# Search in Files – grep

```
grep "hello" file.txt          # Find lines containing "hello"  
grep -i "hello" file.txt      # Case-insensitive search  
grep -e "regex" file.txt       # Match a regular expression  
grep -r "pattern" directory/  # Search recursively in all files  
grep -n "pattern" file.txt     # Show line numbers
```

# Search in Files – grep

```
grep "hello" file.txt          # Find lines containing "hello"  
grep -i "hello" file.txt      # Case-insensitive search  
grep -e "regex" file.txt       # Match a regular expression  
grep -r "pattern" directory/  # Search recursively in all files  
grep -n "pattern" file.txt     # Show line numbers
```

## Example

```
user@linux:~$ grep -n "int" solution.cpp  
3:int main() {  
4:    int n;  
7:    int result = n * 2;
```

Useful for finding specific code, searching logs, or checking patterns across a project.

# Archive – zip / unzip

```
{ zip archive.zip file1 file2 file3      # Zip multiple files  
  zip -r archive.zip directory/        # Zip a directory (requires -r)  
  unzip archive.zip                  # Extract all contents
```

# Archive – `zip` / `unzip`

```
{ zip archive.zip file1 file2 file3          # Zip multiple files  
    zip -r archive.zip directory/           # Zip a directory (requires -r)  
    unzip archive.zip                      # Extract all contents
```

## Example workflow

```
{ zip -r contest.zip CP/Codeforces/Round900/      # Package solutions  
    unzip starter_code.zip                      # Extract downloaded archive
```

`-r` is required when zipping directories (just like `cp` and `rm`).

# Permissions & System Management

chmod, sudo, and apt

# User Permissions

Linux controls file access with three classes and three operations:

## Permission classes

- **Owner** — user who created the file
- **Group** — users in the file's group
- **Others** — everyone else

## Operations

- **r** (read) · **w** (write) · **x** (execute)

## Reading `ls -l` output

```
-rwxr-xr-- 1 user group  script.sh
  ^^^  owner permissions (rwx)
  ^^^  group permissions (r-x)
  ^^^  other perms      (r--)
```

`r` = read, `w` = write, `x` = execute, `-` = denied

# User Permissions

Linux controls file access with three classes and three operations:

## Permission classes

- **Owner** — user who created the file
- **Group** — users in the file's group
- **Others** — everyone else

## Operations

- **r** (read) · **w** (write) · **x** (execute)

## Reading `ls -l` output

```
-rwxr-xr-- 1 user group  script.sh
          ^^^  owner permissions (rwx)
          ^^^  group permissions (r-x)
          ^^^  other perms     (r--)
```

`r` = read, `w` = write, `x` = execute, `-` = denied

If you see "**Permission denied**", you lack the required permission for that operation.

# Change Permissions — `chmod`

```
chmod +x script.sh      # Add execute permission  
chmod -x script.sh      # Remove execute permission  
chmod +r file.txt      # Add read permission  
chmod +w file.txt      # Add write permission  
chmod +rwx file.txt    # Add all permissions
```

# Change Permissions — `chmod`

```
chmod +x script.sh      # Add execute permission  
chmod -x script.sh      # Remove execute permission  
chmod +r file.txt      # Add read permission  
chmod +w file.txt      # Add write permission  
chmod +rwx file.txt    # Add all permissions
```

## When you'll need this

The most common CP use case: making a compiled binary or script executable.

```
g++ solution.cpp -o solution  
chmod +x solution      # Usually not needed – g++ sets this automatically  
.solution  
  
chmod +x compile.sh    # Make a provided shell script runnable  
.compile.sh
```

Only the **file owner** or **root** can change permissions.

# Superuser — sudo

`sudo` runs a command with **root (administrator) privileges**.

```
sudo apt update          # Update package catalogue as root  
sudo apt install g++     # Install a package as root
```

# Superuser — `sudo`

`sudo` runs a command with **root (administrator) privileges**.

```
{ sudo apt update           # Update package catalogue as root  
  sudo apt install g++      # Install a package as root
```

## Why is `sudo` needed?

- System-wide changes require root permissions
- Your normal account is intentionally restricted for safety
- `sudo` temporarily elevates privileges — prompts for your password

The **root user** has unrestricted access to everything. Never log in as root daily — use `sudo` when needed.

# Package Manager — apt

apt is Ubuntu's package manager — an **app store for the terminal**.

```
# Always update the catalogue first  
sudo apt update  
  
# Install packages  
sudo apt install g++  
sudo apt install python3
```

# Package Manager — apt

apt is Ubuntu's package manager — an **app store for the terminal**.

```
# Always update the catalogue first  
sudo apt update
```

```
# Install packages  
sudo apt install g++  
sudo apt install python3
```

```
# Upgrade all installed packages  
sudo apt upgrade
```

```
# Upgrade a specific package  
sudo apt install --only-upgrade g++
```

# Package Manager — apt

apt is Ubuntu's package manager — an **app store for the terminal**.

```
# Always update the catalogue first  
sudo apt update
```

```
# Install packages  
sudo apt install g++  
sudo apt install python3
```

```
# Upgrade all installed packages  
sudo apt upgrade
```

```
# Upgrade a specific package  
sudo apt install --only-upgrade g++
```

```
# Remove a package  
sudo apt remove g++          # Remove the package  
sudo apt purge g++           # Remove package AND config files
```

# C++ Compilation Workflow

g++, flags, warnings, and Makefile

# Installing and Using g++

g++ is the GNU C++ compiler — the standard tool for CP.


```
{ sudo apt update && sudo apt install g++      # Install  
    g++ solution.cpp -o solution                # Compile
```

# Installing and Using g++

g++ is the GNU C++ compiler — the standard tool for CP.

```
{ sudo apt update && sudo apt install g++      # Install  
g++ solution.cpp -o solution                 # Compile
```

What this does



-o solution sets the output name. Without -o, the default is a.out .

# Compilation Flags

Add flags for C++ standard, optimization, and warnings:

```
( g++ solution.cpp -o solution -std=c++17 -O2 -Wall -Wshadow
```

# Compilation Flags

Add flags for C++ standard, optimization, and warnings:

```
{ g++ solution.cpp -o solution -std=c++17 -O2 -Wall -Wshadow
```

Flag	Purpose
<code>-std=c++17</code>	Use C++17 standard (structured bindings, <code>if constexpr</code> , etc.)
<code>-O2</code>	Optimize for execution speed (most judges use this)
<code>-Wall</code>	Enable most compiler warnings
<code>-Wshadow</code>	Warn when a variable shadows another
<code>-Wextra</code>	Even more warnings beyond <code>-Wall</code>

# Compilation Flags

Add flags for C++ standard, optimization, and warnings:

```
{ g++ solution.cpp -o solution -std=c++17 -O2 -Wall -Wshadow
```

Flag	Purpose
<code>-std=c++17</code>	Use C++17 standard (structured bindings, <code>if constexpr</code> , etc.)
<code>-O2</code>	Optimize for execution speed (most judges use this)
<code>-Wall</code>	Enable most compiler warnings
<code>-Wshadow</code>	Warn when a variable shadows another
<code>-Wextra</code>	Even more warnings beyond <code>-Wall</code>

Always compile with `-Wall` during practice — it catches bugs that cost hours in contests.

# -Wall — Warning Examples

## Format string mismatch

```
#include <bits/stdc++.h>
using namespace std;
int main() {
    int x = 5;
    printf("%lld\n", x);
    // Warning: '%lld' expects
    // 'long long int', not 'int'
}
```

## Unused variable

```
#include <bits/stdc++.h>
using namespace std;
int main() {
    int x = 5;
    int y = 9; // Warning: unused 'y'
    printf("%d\n", x);
}
```

# -Wall — Warning Examples

## Format string mismatch

```
#include <bits/stdc++.h>
using namespace std;
int main() {
    int x = 5;
    printf("%lld\n", x);
    // Warning: '%lld' expects
    // 'long long int', not 'int'
}
```

## Unused variable

```
#include <bits/stdc++.h>
using namespace std;
int main() {
    int x = 5;
    int y = 9; // Warning: unused 'y'
    printf("%d\n", x);
}
```

Unused variables often indicate a typo — maybe you meant to use `y` instead of `x`. Warnings catch these before they become bugs.

# -Wshadow — Variable Shadowing

When an inner scope re-declares a variable name, the outer one is **shadowed**:

```
#include <bits/stdc++.h>
using namespace std;
int main() {
    int x = 5;
    for (int i = 0; i < 3; i++) {
        int x = 9;           // Warning: 'x' shadows a previous local
        printf("%d\n", x);   // Prints 9, not 5!
    }
}
```

# `-Wshadow` — Variable Shadowing

When an inner scope re-declares a variable name, the outer one is **shadowed**:

```
#include <bits/stdc++.h>
using namespace std;
int main() {
    int x = 5;
    for (int i = 0; i < 3; i++) {
        int x = 9;           // Warning: 'x' shadows a previous local
        printf("%d\n", x);   // Prints 9, not 5!
    }
}
```

- Common source of bugs in nested loops
- Inner `x` hides outer `x` — compiles silently without this flag
- **Always compile with `-Wshadow`** to catch these early

# Sanitizers – Catch Runtime Bugs

Add sanitizer flags during practice to detect bugs that compile cleanly but crash or produce wrong answers:

```
g++ solution.cpp -o solution -std=c++17 -O2 -Wall -Wshadow \
-fsanitize=address,undefined
```

# Sanitizers – Catch Runtime Bugs

Add sanitizer flags during practice to detect bugs that compile cleanly but crash or produce wrong answers:

```
{ g++ solution.cpp -o solution -std=c++17 -O2 -Wall -Wshadow \
  -fsanitize=address,undefined
```

## What they catch

Sanitizer	Catches
address	Out-of-bounds array access, use-after-free, memory leaks
undefined	Integer overflow, null pointer dereference, shift errors

# Sanitizers — Catch Runtime Bugs

Add sanitizer flags during practice to detect bugs that compile cleanly but crash or produce wrong answers:

```
{ g++ solution.cpp -o solution -std=c++17 -O2 -Wall -Wshadow \
  -fsanitize=address,undefined
```

## What they catch

Sanitizer	Catches
address	Out-of-bounds array access, use-after-free, memory leaks
undefined	Integer overflow, null pointer dereference, shift errors

Sanitizers slow your program ~2-3x — use them for **debugging only**, not for timing. Remove them before measuring performance.

# Running Binaries

After compilation, run with `./`:

```
{ g++ solution.cpp -o solution  
  ./solution }
```

# Running Binaries

After compilation, run with `./` :

```
{ g++ solution.cpp -o solution  
  ./solution
```

Why `./` is needed

Linux looks for executables in `PATH` dirs (`/bin`, `/usr/bin`). Your current directory isn't in `PATH` — `./` tells the shell to look here.

## Measuring execution time

```
{ time ./solution < input.txt  
# real    0m0.032s ← wall clock time (the one you care about)  
# user    0m0.028s ← CPU time in user mode  
# sys     0m0.004s ← CPU time in kernel mode
```

# Running Binaries

After compilation, run with `./` :

```
g++ solution.cpp -o solution  
./solution
```

## Why `./` is needed

Linux looks for executables in `PATH` dirs (`/bin`, `/usr/bin`). Your current directory isn't in `PATH` — `./` tells the shell to look here.

## Measuring execution time

```
time ./solution < input.txt  
# real    0m0.032s ← wall clock time (the one you care about)  
# user    0m0.028s ← CPU time in user mode  
# sys     0m0.004s ← CPU time in kernel mode
```

**CP tip:** If `real` time exceeds ~1-2s for a typical test case, your solution is likely too slow.

# Grader Tasks – Multiple Source Files

Many TFT/IOI problems provide a **grader** – you implement functions, the grader handles I/O.

```
# solution.cpp      - your code (implements functions)
# sample_grader.cpp - provided (contains main, calls your functions)

g++ -o solution solution.cpp sample_grader.cpp -std=c++17 -O2
```

# Grader Tasks – Multiple Source Files

Many TFT/IOI problems provide a **grader** – you implement functions, the grader handles I/O.

```
# solution.cpp      - your code (implements functions)
# sample_grader.cpp - provided (contains main, calls your functions)

g++ -o solution solution.cpp sample_grader.cpp -std=c++17 -O2
```

## Using a provided compile script

```
chmod +x compile.sh && ./compile.sh          # Make executable and run
cat compile.sh                                # Peek at the compilation command
```

You can also compile manually and add any extra flags you need.

# Makefile – Automate Compilation

A **Makefile** saves you from retyping the full `g++` command:

```
% : %.cpp  
    g++ $< -o $@ -std=c++17 -Wall -Wshadow -O2
```

# Makefile – Automate Compilation

A **Makefile** saves you from retyping the full `g++` command:

```
% : %.cpp  
    g++ $< -o $@ -std=c++17 -Wall -Wshadow -O2
```

## How it works

- `%` matches any name → `make solution` looks for `solution.cpp`
- `$<` = source file, `$@` = target name
- **Must use tab** for indentation (not spaces)

```
sudo apt install make      # Install make (if needed)  
make solution             # Compiles solution.cpp → solution  
make A                   # Compiles A.cpp → A
```

Put this Makefile in your working directory – works for any `.cpp` file.

# I/O Redirection & Advanced

Redirects, piping, and command chaining

# I/O Redirection

Redirect input/output from/to files — essential for testing contest problems.

```
{ ./solution < input.txt          # Read stdin from file  
  ./solution > output.txt        # Write stdout to file  
  ./solution < input.txt > output.txt  # Both at once
```

# I/O Redirection

Redirect input/output from/to files — essential for testing contest problems.

```
{ ./solution < input.txt          # Read stdin from file  
      ./solution > output.txt    # Write stdout to file  
      ./solution < input.txt > output.txt # Both at once
```

## Example workflow

```
{ echo "23382338" > input.txt      # Create sample input  
      g++ solution.cpp -o solution   # Compile  
      ./solution < input.txt > output.txt # Run with redirected I/O  
      cat output.txt                 # Check the result
```

# I/O Redirection

Redirect input/output from/to files — essential for testing contest problems.

```
{ ./solution < input.txt          # Read stdin from file  
  ./solution > output.txt        # Write stdout to file  
  ./solution < input.txt > output.txt  # Both at once
```

## Example workflow

```
{ echo "23382338" > input.txt      # Create sample input  
  g++ solution.cpp -o solution       # Compile  
  ./solution < input.txt > output.txt  # Run with redirected I/O  
  cat output.txt                      # Check the result
```

**Why not paste input?** Large inputs cause terminal lag. File redirection is instant and reproducible.

# Why Use I/O Redirection?

# Why Use I/O Redirection?

- **Large inputs** — pasting thousands of lines into the terminal is slow

# Why Use I/O Redirection?

- **Large inputs** — pasting thousands of lines into the terminal is slow
- **Speed** — terminal output is much slower than writing to a file

# Why Use I/O Redirection?

- **Large inputs** — pasting thousands of lines into the terminal is slow
- **Speed** — terminal output is much slower than writing to a file
- **Reproducibility** — keep input files for repeated testing

# Why Use I/O Redirection?

- **Large inputs** — pasting thousands of lines into the terminal is slow
- **Speed** — terminal output is much slower than writing to a file
- **Reproducibility** — keep input files for repeated testing
- **Automated testing** — combine with `diff` to verify output:

# Why Use I/O Redirection?

- **Large inputs** – pasting thousands of lines into the terminal is slow
- **Speed** – terminal output is much slower than writing to a file
- **Reproducibility** – keep input files for repeated testing
- **Automated testing** – combine with `diff` to verify output:

```
{ ./solution < input.txt > my_output.txt  
diff -Z my_output.txt expected_output.txt
```

No output from `diff` = your answer matches exactly.

# Command Chaining – `&&`

Run commands in sequence — next runs only if previous **succeeds**:

```
make solution && ./solution < input.txt          # Compile + run  
sudo apt update && sudo apt install g++           # Update + install  
g++ sol.cpp -o sol && ./sol < in.txt > out.txt && diff -Z out.txt ans.txt
```

# Command Chaining — `&&`

Run commands in sequence — next runs only if previous **succeeds**:

```
make solution && ./solution < input.txt          # Compile + run  
sudo apt update && sudo apt install g++           # Update + install  
g++ sol.cpp -o sol && ./sol < in.txt > out.txt && diff -Z out.txt ans.txt
```

`&&` VS `;`

Operator	Behavior
<code>&amp;&amp;</code>	Next runs <b>only if</b> previous succeeded
<code>;</code>	Next runs <b>regardless</b> of previous result

Using `&&` is safer — if compilation fails, you won't run the old binary.

# Piping — |

The pipe sends **output of one command** as **input to another**:

```
./solution < input.txt | sort          # Sort program output  
./solution < input.txt | wc -l        # Count output lines  
./solution < input.txt | grep "Error" # Search output
```

# Piping — |

The pipe sends **output of one command** as **input to another**:

```
{ ./solution < input.txt | sort          # Sort program output  
  ./solution < input.txt | wc -l        # Count output lines  
  ./solution < input.txt | grep "Error" # Search output
```

Copying output to clipboard

```
{ cat solution.cpp | clip.exe           # WSL (Windows)  
  cat solution.cpp | xclip -selection clipboard # Linux with xclip
```

Useful when output is too long to select with the mouse.

# Terminal Control

## Stopping a command

Press `Ctrl + C` to **terminate** the current command.

Use when:

- Program enters an infinite loop
- Output is flooding the terminal
- You want to cancel an operation

## Useful shortcuts

Shortcut	Action
<code>Ctrl + C</code>	Kill current process
<code>Ctrl + Z</code>	Suspend process
<code>Ctrl + D</code>	Send EOF (end input)
<code>Ctrl + L</code>	Clear screen

# Terminal Control

## Stopping a command

Press `Ctrl + C` to **terminate** the current command.

Use when:

- Program enters an infinite loop
- Output is flooding the terminal
- You want to cancel an operation

## Useful shortcuts

Shortcut	Action
<code>Ctrl + C</code>	Kill current process
<code>Ctrl + Z</code>	Suspend process
<code>Ctrl + D</code>	Send EOF (end input)
<code>Ctrl + L</code>	Clear screen

**CP tip:** If your solution runs for more than a few seconds, it's likely TLE — `Ctrl + C` and optimize.

# Wrapping Up

# Summary

## What we covered

- What Linux is and why it matters for CP
- Navigating the filesystem and terminal
- Essential commands: `ls`, `cd`, `cp`, `mv`, `rm`,  
`cat`, `mkdir`, `echo`, `diff`, `grep`, `zip`
- Permissions with `chmod` and `sudo`
- Package management with `apt`
- C++ compilation with `g++` and flags
- I/O redirection, piping, and chaining

## Compilation cheat sheet

```
# Compile  
g++ sol.cpp -o sol -std=c++17 \  
-O2 -Wall -Wshadow  
  
# Run with I/O redirection  
../sol < input.txt > output.txt  
  
# Check output  
diff -Z output.txt expected.txt  
  
# All in one  
make sol && ./sol < in.txt \  
> out.txt && diff -Z out.txt ans.txt
```

# Try It Yourself

Practice everything you've learned in the interactive Linux Playground:

The screenshot shows the Linux Playground interface. At the top left is the title '>\_ Linux Playground' with a refresh icon. To its right is a progress bar labeled 'Progress' and '0/6 Levels Complete'. Below this is a mission panel titled 'L1 Project Scaffolding' with a completion count of '0/5'. The panel contains descriptive text about organizing a workspace and creating a nested directory structure. A mission step '1. Create the nested path' is listed, along with a hint button. To the right of the mission panel is a large text area containing a welcome message and instructions, followed by a terminal prompt 'user@linux:~\$'.

Welcome to Linux Playground – Competitive Programming Edition

Type "help" for available commands. Complete the missions on the left panel.

user@linux:~\$

Complete all 6 mission levels — from basic navigation to a full contest simulation workflow.

# Questions?

linux-playground.vercel.app